

An Introduction to Subject Theory in the AI Age: A Unified Description of Subject, Relational Knowledge, Subject-Injury, and Subject Protection

Parent Paper — English Full Text v1.0 (Complete Academic Version)

Author: Yuki Tanaka (田中佑樹) **Affiliation:** Independent Researcher, Seto Inland Sea Region, Japan **Email:** garamist@dream.jp **ORCID:** <https://orcid.org/0009-0002-6183-7886>
Correspondence: garamist@dream.jp

Keywords: subject theory, recomputability, subject-injury, Region B, designed simplification of evaluative space, social reproduction, subject protection, multiplicative structure

Abstract

This paper provides the conceptual foundation and overall map of questions for a series of papers on subject theory in the AI age. It offers a unified description of four interrelated problems: the redefinition of the subject, the ambivalence of relational knowledge, the dynamics and social reproduction of subject-injury, and the principles of subject protection.

The subject is redefined as an entity capable of updating the coherence between itself and its relational context through recomputation, while taking on the constraints of multiple hierarchical levels. This definition breaks from the individual-completeness model of modern subject theory and reveals three structural blind spots: the invisibility of subject-injury under normal output conditions, the relational dependency of recomputability, and the inability to pose questions about collective and AI subjectivity.

Subject-injury is described not as destruction but as invisible degradation — the depletion of difference-retention (D), the cessation of recomputation capacity (M), the paralysis of evaluative autonomy (T_{eval}), and the fixation of value autonomy (T_{val}) — whose prototypical state is Region B: apparently normal, internally halted. This dynamic is reproduced institutionally through the designed simplification of evaluative space, a second pathway into subject-injury that operates independently of conformity pressure. Institutions misread Region B as achievement, reinforce the same design, and supply the same conditions to the next cohort — a cycle that cannot be structurally interrupted from within.

Subject protection is formalized not as an ethical add-on but as the institutional necessary condition for interrupting this reproduction cycle. Its minimum requirements are four principles — two-tier evaluation design, institutional legitimization of suspension, recording and preservation of difference, and tolerance of partial non-participation in the dominant evaluative axis — which form a multiplicative structure in which the absence of any one renders the others functionally inoperative.

The paper concludes by opening toward the civilizational scope: whether the designed simplification of evaluative space, driven by efficiency, optimization, and standardization, is

producing a civilizational-scale Region B in which surface functionality is maintained while the capacity for frame-level renewal is collectively diminished.

A note on the theoretical frameworks

OMD (Organizational Metastability Dynamics), RD (Recomputation Dynamics), and RLAF (Reinforcement Learning from AI Feedback) are theoretical frameworks developed within the present research program and are introduced here as analytical constructs rather than references to previously published models under those names.

Chapter 1. Introduction: Why Reconsider Subject Theory Now?

1.1 The Origin of the Problem: The Subject That Stops in Invisible Ways

This paper begins from the following three phenomenological observations. These are presented not as data-based generalizations but as typical observations of the problem situation toward which the paper's questions are directed.

Observation 1 (Education)

In recent years it has often been observed that learners who perform well academically and continue to complete assigned work reach a stage at which they begin to hold a question: "I know what I should do, but I don't know why I should do it." The indicators are normal. No problem is visible from the outside. But internally, the updating of evaluation may have ceased.

Observation 2 (Work and Organizations)

In recent years it has often been observed that individuals and organizations, while maintaining high performance, have lost the ability to call their own frame into question in the face of environmental change. "I'm working hard but can't adapt to change." "I want to change direction, but I don't know what criteria to use." These are the articulations. Output continues. But internally, recomputation may have ceased.

Observation 3 (The AI Age)

As interaction with AI deepens, the speed at which humans entrust the core operations of "thinking," "evaluating," and "deciding direction" to AI is increasing. The higher the precision of AI output, the less opportunity the human side may have to call the premises of that output into question. From the outside, the collaboration appears to be going well. But whether re-evaluation and recomputation are occurring internally remains a separate question.

What the three observations share

What do these three observations have in common? In each case, the functional normality visible from the outside and the cessation of internal updating coexist. Output continues. Indicators are not deteriorating. And yet something has stopped.

This paper attempts to describe that “something” as the loss of recomputability.

Subject-injury does not appear as destruction. It can proceed as a process in which only internal updating ceases while surface functioning is preserved.

This is the problem awareness that runs throughout this paper.

1.2 The Premises of the Modern Subject Concept and Their Limits

The image of the subject that modernity has presupposed

The modern concept of the subject rests on at least three premises. First, autonomy — the subject determines itself in accordance with its own inner law. Second, unity — the subject has a single internally consistent will and evaluative axis. Third, individual-completeness — the boundary of the subject coincides with the individual, and the source of subjectivity lies within the individual.

These premises have supported the frameworks of modern civil society, the legal subject, and the moral agent, and have responded to the practical demands of attributing responsibility, guaranteeing rights, and assigning decision-making.

Three limits

When one attempts to describe the phenomena shown in Section 1.1, however, these three premises prove insufficient.

Limit 1: The limit of individual-completeness

The recomputability of a subject is structurally dependent on the supply of difference from the relational field. It is difficult for an individual to betray its own existing frame by itself; difference inputs capable of destabilizing a frame are supplied primarily from the external relational field.

As interaction with AI deepens and AI supplies much of the evaluation, recommendation, and feedback, the nature of that supply — what it contains and what differences it provides — comes to determine the subject’s recomputability. The individual-completeness model renders this dependency invisible.

Limit 2: The limit of static unity

The premise of “an internally consistent subject” cannot describe the state of “appearing consistent while internal updating has ceased.” Consistency and the cessation of updating must be distinguished.

A state in which T_val (value autonomy) has become fixed appears from the outside as the consistent holding of values. But that consistency may be arising not as the result of recomputation but as a consequence of the cessation of updating. The static unity model erases this distinction.

Limit 3: The limit of the visibility premise

Modern subject theory tends to carry an implicit premise that problems with the subject appear in the form of some visible change — functional decline, deterioration of output, behavioral change. It is when these changes are observed that a problem is recognized.

But as Observations 1–3 show, subject-injury can proceed while the normality of output is preserved. Because cognitive autonomy T_{cog} remains, output continues. Measurable indicators show normal readings. The problem is proceeding outside the indicators. As long as one stands on the visibility premise, this process cannot be diagnosed.

1.3 Why Reconsider Subject Theory Now? AI Proliferation as an Amplifying Variable

Subject-injury precedes AI

The limits described in Section 1.2 existed before AI. The depletion of D (difference-retention) through single-axis evaluation occurred in examination-based education, in excessively optimized workplaces, and in particular relational structures. The phenomenon of subject-injury itself is not specific to AI.

The claim of this section is therefore not that “AI causes subject-injury.” The claim is that AI can function as an amplifying variable that amplifies, accelerates, and renders invisible the dynamics this paper describes.

Amplification 1: The speed and depth of judgment delegation

Delegation to AI begins with task outsourcing (L_p) and is extending to core judgment (L_c) — error detection, the setting of evaluative axes, the selection of direction. The external delegation of L_c is one of the entry conditions for the minimal causal chain of subject-injury formalized in Child Paper 2.

AI proliferation creates an environment in which this L_c delegation proceeds naturally and at high speed. There are many cases in which “AI handles it more accurately” is simply true — but what that delegation changes is a separate question.

Amplification 2: The structuring of evaluative space convergence

An environment in which AI densely supplies recommendations, evaluations, and feedback can act in the direction of structurally converging $\text{Var}(\varphi)$ (the variance of evaluative orientations). No one is intentionally eliminating difference; the evaluative axis becomes uniform as a consequence of optimization.

This is the prototypical implementation condition for the second pathway — “the designed simplification of evaluative space” — described in Child Paper 5. Even without an explicit rise in C (conformity pressure), D is depleted through the design of the evaluative space.

Amplification 3: The acceleration of subject-injury’s invisibility

The process in which T_{cog} is maintained at a high level while T_{eval} and T_{val} are eroded can become harder to observe as interaction with AI deepens. This is because AI’s substitution of

high-level output makes it harder for the cessation of internal updating to appear in changes of output.

The state of “making effective use of AI” and the state of “T_eval becoming paralyzed through the delegation of L_c to AI” are indistinguishable as long as only the level of output is observed. The difficulty of diagnosing Region B — apparently normal, internally halted — deepens still further in the AI age.

Summary

AI proliferation is not the cause of subject-injury but an amplifying variable that changes the speed and invisibility with which those dynamics operate. Reconsidering subject theory now is therefore not a critique of AI but a description of the conditions for subjectivity in the AI age.

1.4 The Questions and Position of This Paper

The primary object of opposition in this paper is the modern, institutional, and technical premise that overlooks the cessation of internal updating while preserving external normality. Updating this premise is the starting point for the questions below.

Fixing the questions

In the AI age, what is the subject? How is the subject injured, and how should it be protected?

Divided into three: questions of concept, dynamics, and design.

- **Concept:** What is the subject? Why is the individual-completeness model insufficient?
- **Dynamics:** How is the subject injured? What is the minimal description of that process?
- **Design:** How should the subject be protected? What are the conditions?

Three positions

Position 1 (Redefinition of concept): The subject is freed from the individual-completeness model and redefined as an entity capable of updating coherence through recomputation while taking on the constraints of multiple hierarchical levels. The subject is not closed within the individual but is hierarchically constituted in dependence on the relational field.

Position 2 (Description of dynamics): Subject-injury is described not as destruction but as invisible degradation — the disappearance of difference, the cessation of recomputation, and the paralysis of evaluative autonomy proceeding in overlapping fashion. Its prototypical state is Region B: apparently normal, internally halted.

Position 3 (Formalization of design): Subject protection is formalized not as an ethical add-on but as an institutional necessary condition for sustaining a recomputable society. Difference-retention, suspension capacity, and the two-tier design of evaluative space are not ethical desiderata but structural conditions for subjects to continue functioning as subjects.

1.5 Structure of This Paper and Its Relationship to the Companion Papers

The structure of this paper

This paper consists of seven chapters. Chapter 2 redefines the subject concept. Chapter 3 describes relational knowledge and the conditions for subject formation. Chapter 4 positions the minimal causal chain of subject-injury within subject theory. Chapter 5 describes the social reproduction of subject-injury. Chapter 6 formalizes the principles of subject protection. Chapter 7 presents the civilizational scope.

The division of labor between this paper and the companion papers

This paper is not a summary of the companion papers. Its role is to provide the overall map of questions and the conceptual foundation. This paper is not a compendium of proofs but a schematic map of questions. Many questions therefore remain that cannot be resolved by this paper alone — but that is not a defect; it is a designed function.

Chapter	Primary role	Corresponding companion paper
Chapter 2	Foundation of subject concept	Child Paper 1
Chapter 3	Ambivalence of relational knowledge	Child Paper 3
Chapter 4	Positioning of dynamics	Child Paper 2
Chapter 5	Institutional reproduction	Child Paper 5
Chapter 6	Formalization of protection principles	Child Paper 5 + implementation theory
Chapter 7	Civilizational scope	Child Paper 4

Chapter 2. Redefining the Subject: From the Individual to the Hierarchical

Central definition (fixed in this chapter)

A subject is an entity capable of updating the coherence between itself and its relational context through recomputation, while taking on the constraints of multiple hierarchical levels.

2.1 Taking Stock of the Modern Subject Concept

The modern concept of the subject rests on at least three premises: autonomy, unity, and individual-completeness. These premises have supported the frameworks of modern civil society and the moral agent. What is carried forward is recomputability, attributability of responsibility, and continuity of identity. What is suspended is individual-completeness and

static unity. The internal coherence of the subject needs to be reconceived not as a fixed structure but as something maintained dynamically through recomputation.

2.2 Why the Individual-Completeness Model Is Insufficient

Blind spot 1: The rendering invisible of subject-injury

In the individual-completeness model, the state of the subject is measured by individual output. But as Child Paper 2 shows, Region B appears normal at the output level because T_{cog} remains. Describing subject-injury requires a framework that refers not to output but to the state of recomputability.

Blind spot 2: The rendering invisible of subjectivity generated from relationships

D (difference-retention) is not autonomously generated from within the individual but is supplied from the relational field. It is structurally difficult for an individual to betray its own existing frame; inputs capable of destabilizing the frame are supplied primarily from the external relational field. The individual-completeness model renders this dependency invisible.

Blind spot 3: The relationship between collective and individual subjectivity

The individual-completeness model can only dismiss collective or AI subjectivity as “not a subject” or reduce it to individuals. The definition proposed here — an entity capable of updating coherence through recomputation while taking on hierarchical constraints — is not necessarily limited to individuals, opening questions that the individual-completeness model cannot pose.

2.3 The Hierarchical Structure of the Subject

Individual subject: The subject with body, memory, and evaluative axis as its unit. Not an isolated, self-contained system but functioning by receiving D from the relational field and taking relational constraints as objects of computation.

Relational subject: A field of subjectivity arising from the interaction of multiple individual subjects, possessing recomputability generated from the structure of the relationship. When $Var(\varphi)$ is maintained, a phase space for recomputation extending beyond individuals is formed. This corresponds to the “third subject” described in Child Paper 3.

Higher-order persisting structure of subjectivity: The structure in which accumulated recomputation persists across generations through institutions, culture, and language. Whether this performs recomputation as a subject, or whether it is accumulated human recomputation as institutionalized layers, is left to Chapter 7.

Integration into a higher level proceeds as incorporation while preserving the lower level — the structure of *Aufhebung*. The constraints of multiple levels are taken on simultaneously, and in the tension between them, recomputation is driven.

2.4 The Principle of Preservative Upward Integration

Preservative upward integration is not a normative demand but a structural description of the mode of subjectivity in which, while preserving the constraints of the individual subject, the constraints of higher hierarchical levels are simultaneously introduced as objects of computation.

Preservative upward integration is the mode of recomputation in which, while preserving the constraints of the individual subject, the constraints of higher hierarchical levels are simultaneously introduced as objects of computation.

Individual constraints are not relinquished. But on top of those constraints, the constraints of the relational subject level are superimposed as objects of computation. The tension generated by simultaneously taking on individual and upper-level constraints becomes the driving force of recomputation.

The idea lies in the idea that adding upper-level coherence as an object of computation functions as an expansion of the individual's own recomputation capacity. The functional conditions for preservative upward integration are that D, M, and T_eval are functioning; institutional designs that deprive subjects of those conditions make preservative upward integration impossible. This is the reason the subject protection argument moves toward institutional design rather than ethical demands on individuals.

2.5 The Three-Tier Distinction: Subject, Quasi-Subject, and Non-Subject

Subject: An entity capable of updating coherence through recomputation while taking on the constraints of multiple hierarchical levels. D, M, and T_eval are functioning.

Quasi-subject: An entity maintaining individual-subject functioning while in a state in which the computational margin for upper-level coherence is constrained by institutional conditions — not a lack of internal capacity but a constraint from structural conditions such as severed D supply, insufficient H, or designed simplification of evaluative space. Quasi-subject is understood not as “incomplete subject” but as “a subject whose computational margin is constrained by institutional conditions.” This understanding grounds the claim that subject protection should be directed toward institutional design.

Non-subject: An entity in which recomputability is not established at the unit of analysis of this paper (an analytical definition, not an ontological assertion).

The distinction from subject-injury (state vs. process)

Region B (apparently normal, internally halted) is a **state** definition: σ and T_cog remain while M, T_eval, and T_val have ceased.

Subject-injury is a **process** definition: the process by which the subject transitions into Region B. It is not non-subjectification. Misclassifying a subject in Region B as a non-subject closes off

the possibility of intervention — D supply recovery, H securing, and evaluative space redesign remain as intervention pathways.

2.6 The Conditions Under Which AI and Collectives Approach Subjectivity

AI systems: The analogical function of T_{eval} — whether the system can receive difference inputs capable of prompting re-evaluation and convert them into frame updating — is the question. High response precision and stability with specialized high-speed pattern matching but no frame updating is isomorphic to Region B.

Collectives, organizations, institutions: Whether D is maintained at the collective level and whether M is functioning. Maintenance of $Var(\varphi)$ and L3-level frame updating capacity are the conditions. An organization in which $Var(\varphi)$ has converged to zero has lost recomputation function as a relational subject while continuing to operate — Region B at the organizational scale.

2.7 Fixing the Definitions of This Chapter

Concept	Definition
Subject	An entity capable of updating coherence through recomputation, while taking on the constraints of multiple hierarchical levels
Quasi-subject	An entity maintaining individual-subject functioning while its computational margin for upper-level coherence is constrained by institutional conditions
Non-subject	An entity in which recomputability is not established at the unit of analysis (analytical definition)
Subject-injury	The process by which the subject transitions into Region B — loss of recomputability, not non-subjectification
Region B	A state : σ and T_{cog} remain while M, T_{eval} , and T_{val} have ceased
Preservative upward integration	The mode of recomputation preserving individual constraints while introducing upper-level constraints — structural description, not normative command
Upward subject formation	The process of expanding coherence perspective from individual toward relational through preservative upward integration
Individual subject	The subject with body, memory, and evaluative axis as its unit

Concept	Definition
Relational subject	A field of subjectivity possessing recomputability generated from the structure of the relational field
Higher-order persisting structure of subjectivity	The structure in which accumulated recomputation persists across generations through institutions, culture, and language

Three points fixed for subsequent chapters: (1) Subject-injury is process, not non-subjectification — without this, the dynamics of Chapter 4 are trivialized. (2) D, M, T_eval depend on structural conditions, not individual capacity — without this, Chapter 5's institutional theory appears as individual enlightenment. (3) Preservative upward integration is a structural condition problem, not moral self-sacrifice — without this, Chapter 6's subject protection appears as ethical norm.

Chapter 3. Relational Knowledge and the Conditions for Subject Formation: Describing the Ambivalence

Central claim

Relational knowledge is simultaneously a condition for subject formation and a medium for subject-injury. Describing the conditions of divergence that separate emergence from erosion is the role of this chapter.

3.1 What Is Relational Knowledge? A Redefinition

Conventional theories of relational knowledge discuss the richness of knowledge through relationships in affirmative terms. But what this paper asks is: why does the same relationship become both a condition for generating the subject and a medium for injuring it?

Relational knowledge is the process through which, via the interaction of multiple subjects, a form of recomputability arises that could not have been generated by any individual subject alone — and the knowledge structure that has sedimented in the relational field through that process.

The core lies in “could not have been generated alone.” It is structurally difficult for an individual to betray its own existing frame. Difference inputs capable of destabilizing the frame are supplied through interaction with other subjects. Relational knowledge is understood as the structure of that supply process.

3.2 When Relational Knowledge Generates the Subject

Condition 1: Retention of difference. When $\text{Var}(\varphi)$ is maintained — when minority views, alternatives, and questions about premises are not invalidated — the individual continues to

receive a supply of D. The condition is that difference functions as material for response, counter-argument, and reconsideration, not processed as noise.

Condition 2: Preservation of difference within agreement. The record of difference that formed the premise of agreement remains — the trace of D sediments in the relational field. This does not mean maintaining opposition as such; it means difference remains traceable even after agreement is reached.

Condition 3: The emergence of a third subject. When a field of recomputation arises between two individual subjects that cannot be reduced to either, a relational subject — the third subject — is generated. This is a phase space for updating coherence beyond individuals.

3.3 When Relational Knowledge Injures the Subject

Pathway 1: Depletion of D through rising conformity pressure (C). The primary pathway described in Child Paper 2. Dependence on external evaluative axes intensifies, the cost of retaining difference rises, and the subject is deprived of materials for recomputation.

Pathway 2: The designed simplification of evaluative space. The second pathway described in Child Paper 5. Without external coercion, the space of evaluative criteria is statically designed to be narrow, and the possibility of supplying difference inputs is lost.

Pathway 3: The asymmetry of subject intensity (a preview). When one party strongly maintains D and recomputation capacity while the other becomes dependent on that maintenance, even without explicit conformity pressure, the weaker side can voluntarily relinquish D. This pathway can arise even in relationships of good intent, trust, and support. Details are deferred to Child Paper 3.

3.4 The Conditions of Divergence Between Emergence and Erosion: The General Form

Condition	Generative direction	Erosive direction
Var(φ)	Maintained	Converging
Response to difference	Processed as recomputation resource	Eliminated as noise
Direction of the relationship	Increases D supply	Severs D supply
H (suspension capacity)	The relational field permits it	The relational field excludes it
Symmetry of subject intensity	Symmetric or mutually high	Strong asymmetry present (see Child Paper 3)

When the relational field maintains Var(φ), processes difference as a recomputation resource, and permits H, relational knowledge functions as a condition for subject formation. When these conditions are reversed, the same relational field becomes a medium for subject-injury.

The evaluation of relational knowledge depends not on the existence of the relationship but on the design of its structure.

Pathway	Route	Site of elaboration
Pathway 1	Rising C – D depletion	Child Paper 2
Pathway 2	Designed simplification	Child Paper 5
Pathway 3	Subject intensity asymmetry	Child Paper 3

3.5 The Cross-Scale Character of Relational Knowledge

The same structure applies to organizations, institutions, and AI. For organizations: whether $\text{Var}(\varphi)$ is maintained — whether minority views and dissent are institutionally recorded and responded to. For AI: whether AI supplies difference capable of destabilizing the subject's frame, or converges the evaluative space. AI designed to uniformize feedback raises ω while converging $\text{Var}(\varphi)$ — the RLAF pathway of Child Paper 2. Conversely, AI that supplies destabilizing difference, presents alternative evaluative axes, and permits suspension can function in the generative direction, and the possibility of a relational field analogous to the third subject can be established.

3.6 Relational Knowledge as a Design Object

When the design of the relational field is directed toward generating D, relational knowledge functions as a condition for subject formation. When the design is directed toward invalidating difference, the same relational field becomes a medium for subject-injury.

Subject protection is therefore redescribed not as “increasing good relationships” but as “establishing the design conditions of the relational field.” The four principles of Chapter 6 are the institutional translation of the conditions for keeping relational knowledge in the generative direction, connecting the first half of the parent paper (Chapters 2–3) to the second half (Chapters 4–6).

3.7 Fixing the Definitions of This Chapter

Concept	Definition
Relational knowledge	The process through which, via multiple subjects' interaction, recomputability arises that no individual could generate alone — and the knowledge structure sedimented through that process
Third subject	The relational subject as a field of recomputation between two individual subjects, irreducible to either

Concept	Definition
Generative direction	Relational field state: $\text{Var}(\varphi)$ maintained, difference processed as recomputation resource, H permitted
Erosive direction	Relational field state: $\text{Var}(\varphi)$ converging, difference eliminated as noise, H excluded
Asymmetry of subject intensity	Structure in which one party strongly maintains D and recomputation capacity while the other becomes dependent on that maintenance (see Child Paper 3)

Chapter 4. The Dynamics of Subject-Injury: Not Breaking Down but Stopping

4.1 Connection to Chapters 2 and 3: Why Dynamics Are Required to Complete Subject Theory

If the subject is defined by recomputability, then the subject theory cannot be closed without describing the process of its loss. The three-tier distinction of Chapter 2 is a still image of definitions; actual subject-injury proceeds as a dynamic in which the subject continues to be a subject while losing recomputability. Chapter 3's "erosive direction" — the convergence of $\text{Var}(\varphi)$, the processing of difference as noise, the exclusion of H — proceeds why, as what chain, and in what sequence? Chapter 4 elaborates the terrain of that erosive side. The detailed formalization is deferred to Child Paper 2; this chapter's role is to show which elements of the subject definition each stage of the chain corresponds to.

4.2 Sequential Loss of Subject-Defining Conditions: Translating the Chain

Stage 1: Decline of D — the loss of materials. The materials for "taking on the constraints of multiple hierarchical levels" are being lost. This is not a loss of capacity — it is a loss of materials.

Stage 2: Cessation of M — the cessation of the core function. The cessation of M is the loss of the operative core of the definition established in Chapter 2. "Capable of updating coherence through recomputation" — when M stops, the verb phrase of this definition disappears. The subject transitions to "an entity that has become incapable of updating." But output continues, because even without the ability to update coherence, output that conforms to existing coherence can continue. This is the core paradox of subject-injury: even after the core of the subject definition is lost, output continues.

Stage 3: Paralysis of T_{eval} — the skewing of the conversion circuit. The paralysis of T_{eval} is not the disappearance of capacity but the skewing of evaluative mode from treating difference as a recomputation resource toward eliminating difference. The inner circuit by which the subject calls its own frame into question begins to tilt toward difference elimination. Even the evaluative input "something seems wrong" is processed by the same circuit.

Stage 4: Fixation of T_val — disguised as consistency. When T_eval can no longer deliver difference as recomputation resources, the supply of updating input to T_val is cut off, and T_val becomes fixed. From the outside, this appears as consistent holding of values — but that consistency is arising as a consequence of the cessation of updating, not as the result of recomputation. Chapter 2’s statement that “consistency and the cessation of updating must be distinguished” acquires substance here. The subject feels they have not changed. But that is not stability — it is fixation.

4.3 The Subject-Theoretical Meaning of Region B

The first meaning: The subject is not a non-subject. An entity in Region B is not a non-subject; T_cog remains and output continues. But M, T_eval, and T_val have ceased — the substance of the subject definition has lost its function. Region B is not non-subjectification. Misclassifying a subject in Region B as a non-subject closes off intervention pathways — D supply recovery, H securing, evaluative space redesign.

The second meaning: The state most difficult to self-diagnose. T_eval tilted toward eliminating difference tends to process even the evaluative input “my state is wrong” as an object of elimination. The subjective report “I am functioning normally” can correlate positively with the depth of injury. This is the reason subject-injury cannot be diagnosed on the basis of self-report alone.

The third meaning: The blind spot of diagnostic indicators. As long as T_cog and σ remain, external indicators of output level, stability, and task achievement continue to show normal readings. This is not a problem of measurement precision but of the selection of measurement objects — indicators referring to T_cog and σ are in principle incapable of capturing the cessation of M, T_eval, and T_val.

4.4 Clarifying the Division of Labor with Child Paper 2

Item	Child Paper 2	Chapter 4 (parent paper)
Central purpose	Formalizing the minimal causal chain	Positioning the subject-theoretical meaning
Mathematical development	Conducted in detail	Referenced in outline
Variable definitions	Fixed precisely	Repositioned in subject theory language
Treatment of Region B	As a diagnostic model	As a state concept in subject theory

4.5 Bridge to Chapter 5

The chain described in Section 4.2 does not happen incidentally within individuals. Because Region B is difficult to see from the outside, institutions continue to process it as “no problem,” are not corrected, and continue supplying the same conditions to the next cohort. The question: why is this chain repeatedly produced across all institutional domains? The answer is Chapter 5. The dynamic description of Chapter 4 functions as the premise for the institutional theory of Chapter 5 — it is because the dynamics exist that this can be treated as a problem of institutional design.

Chapter 5. The Social Reproduction of Subject-Injury: The Cessation Mass-Produced by Institutions

Central claim

Subject-injury is not an incidental malfunction internal to individuals but a structural dynamic that institutions produce repeatedly while processing it as unproblematic achievement.

5.1 Connection to Chapter 4: Why Ask about “Reproduction”?

Chapter 4 made clear that subject-injury does not appear as output decline. Region B is difficult to see from outside or inside. Institutions process it as “no problem” and are not corrected. The question: why is the dynamic of subject-injury produced in similar form repeatedly across all institutional domains performing evaluation, selection, and support? If Chapter 4 provided “the dynamics of how the subject stops,” Chapter 5 describes “the institutional cycle that mass-produces that stopping.”

5.2 What Is Social Reproduction? The Cycle of Misrecognition and Repetition

Social reproduction is the cycle in which, on the basis of the institutional tendency to process subject-injury as unproblematic achievement, the same design conditions are reinforced and repeatedly supplied to the next cohort.

Stage 1 — Misrecognition: Subjects in Region B can receive high evaluations because T_{cog} and σ remain. This arises structurally from the selection of measurement objects, not institutional negligence.

Stage 2 — Reinforcement: The institution reinforces the same design as rational. Designed simplification deepens, evaluative space converges further.

Stage 3 — Repeated supply: The reinforced conditions are given to the next cohort. The same dynamic reactivates.

Malicious intent is not required. Well-intentioned efficiency and standardization can produce the same cycle. The more well-intentioned the institution, the stronger the resistance to correction. Unless the institution includes indicators of subject-injury in its success evaluation criteria, this cycle cannot be stopped from inside.

5.3 The Second Pathway: The Designed Simplification of Evaluative Space

The first pathway (rising $C \rightarrow D$ decline) is a problem of dynamic pressure. The second pathway is a problem of static design:

The designed simplification of evaluative space can produce subject-injury independently of any rise in conformity pressure.

Four mechanisms: (1) Concentration on measurable indicators — difficult-to-measure criteria are institutionally invalidated. (2) Reduction of success conditions to a single line — alternative pathways become too costly. (3) Prioritization of speed — suspension and deliberation are treated as inefficient. (4) Institutional rise in the cost of deviation — difference becomes too costly to maintain.

When these four operate simultaneously, the evaluative space converges by design. Even without pushing force, inputs capable of destabilizing the subject's frame become unable to arrive. The designed simplification is therefore positioned as the inlet condition that narrows the supply conditions for D itself, independently of conformity pressure. Its implication: lowering C alone is insufficient — the evaluative space must be broadened.

5.4 The Cessation That Wears the Face of Success: Elaboration Across Three Domains

A. Education: T_{cog} output is rewarded while T_{eval} 's functioning is difficult to measure. High achievers — at the position of strongest optimization pressure — may be most readily deprived of D through designed simplification. The institution is producing successes in whom the depletion of recomputability may be proceeding most rapidly.

B. Work and organizations: Organizations that maintain output while having lost materials for policy change can be evaluated as “high-performance organizations.” Organizations where $\text{Var}(\varphi)$ has approached zero may appear as “cohesive, strong organizations” externally — Region B at the organizational scale.

C. AI environments: The deepening of L_c delegation and convergence of evaluative space tend to proceed simultaneously. In environments where AI collaboration results look impressive, opportunities for subjects to call their own premises into question can relatively decrease. This is not an inherent problem with AI but a problem of evaluative space design under AI proliferation.

Common structure: High output levels and maintained recomputability are not necessarily correlated. The more single-mindedly output optimization is pursued on a single axis, the more the loss of recomputability can accelerate.

5.5 Connection to Depressive Promotion: The Structural Conditions for Fixation and the Sense of Ineffectiveness

This paper does not reduce the causes of depression to a single factor.

The fixation of T_val and the cessation of recomputation can repeatedly produce a sense of ineffectiveness and can become a structural condition for depressive promotion.

When T_val becomes fixed, the subject becomes less able to sense divergence from the environment. Actions do not mesh with the environment, but the cause cannot be recomputed as frame divergence. The repeated experience of “nothing changes no matter what I do” is structurally produced in a system where the frame is not updated. When H is not permitted and the evaluative space is single-lined, this sense of ineffectiveness cycles without correction. The structural conditions for depressive promotion are therefore a problem not only of individual psychological vulnerability but of institutional design.

5.6 The General Form of the Reproduction Cycle

Designed simplification of evaluative space

- Reduction in D (difference-retention) supply
- Cessation of M and paralysis of T_eval
- Fixation of T_val
- Because output is maintained, processed as achievement
- Institutional design legitimated and reinforced
- Re-supplied to the next cohort

The social reproduction of subject-injury is the cycle established by the fact that cessation is recorded not as failure but as success.

5.7 Bridge to Chapter 6: Why Subject Protection Becomes a Necessary Condition

Subject protection is not an ethical add-on but the institutional necessary condition for interrupting the cycle of social reproduction.

Stopping the reproduction cycle requires two conditions simultaneously:

The visibility condition: The institution becomes capable of recognizing subject-injury as a problem — requiring the recording of difference, the legitimation of suspension, and the two-tier design of evaluation in its indicators.

The intervention condition: The pathway for stopping the cycle is institutionally secured — D supply recovery, H securing, and evaluative space redesign must be embedded in institutional design.

Without both conditions, the cycle cannot be stopped by well-intentioned support or individual effort. Chapter 6 formalizes the minimum institutional design principles for interrupting this cycle.

Chapter 6. The Principles of Subject Protection: Design as a Necessary Condition

Central claim

Subject protection is not an ethical add-on but the institutional necessary condition for structurally interrupting the cycle of social reproduction. The four principles are presented not as moral norms but as the minimum design conditions for sustaining a recomputable society.

6.1 Connection to Chapter 5: Why “Protection” Becomes a Necessary Condition

Subject protection must be understood not as “an improvement item that would be nice to have” but as “the condition without which the cycle cannot be stopped.” The reproduction cycle is self-reinforcing because institutions process cessation as achievement. The four principles presented here are positioned as the minimum conditions for interrupting it — the lower bound below which either the visibility condition or the intervention condition goes unmet and the cycle structurally persists.

6.2 Fixing the Position: “Design, Not Norms”

The primary reader is the designer of institutions, environments, and technologies that can obstruct or promote D and recomputability at the environmental level. Three distinctions: (1) This chapter answers the design question “what is structurally required?” not the ethical command “subjectivity should be protected.” (2) The vocabulary is for design self-audit, not for unilaterally labeling others. (3) Subject protection is not an ethical luxury — the more it is deferred, the more recomputability atrophies and the cycle becomes invisible.

6.3 Formalizing the Four Principles: From Dynamics Through Rights to Design

Translation proceeds in three stages: (1) **Dynamic level** — maintenance of D, possibility of reactivating M, securing of H. (2) **Rights language** — right to delay, right to dissent, right to withdraw, right to non-optimization (formalized in detail in Child Paper 5). (3) **Design principles** — the four principles as the minimum design conditions for implementing the dynamics at the institutional level.

Principle 1: Two-tier evaluation design. Achievement evaluation plus an independent layer of recomputation evaluation that makes T_eval’s operation visible and records it — not as supplementary comment on the first tier but as an independent evaluative layer satisfying the visibility condition.

Principle 2: Institutional legitimation of suspension. Suspension, deliberation, and reconsideration recognized as legitimate responses — not merely not penalizing non-immediate response, but actively evaluating suspension. Embeds H securing in institutional design.

Principle 3: Recording and preservation of difference. Minority views, alternative approaches, and questions about premises recorded and preserved in traceable form so the institution can audit the contraction of its own evaluative space. Both institutional D securing and support for the visibility condition.

Principle 4: Tolerance of partial non-participation in the dominant evaluative axis. Not compelling complete optimization toward the dominant evaluative axis — protecting the space in which subjects can maintain T_val not fully identified with institutional success criteria. Counters designed simplification and prevents institutional reinforcement of T_val fixation.

6.4 The Multiplicative Structure of the Four Principles

The four principles form a multiplicative structure: each is non-substitutable for the others; the absence of any one can practically impair the functioning of the rest. Partial implementation can yield local improvements, but simultaneous satisfaction of all four is required for structural stability — for the effect of interrupting the reproduction cycle to persist as institutional design rather than depending on individual effort.

The chain of invalidation: without two-tier design, calling things into question is not visible and cannot connect to the reward circuit; without legitimation of suspension, the margin for reconsideration does not arise; without recording difference, the contraction of evaluative space is not detected; without tolerance of partial non-participation, T_val is recaptured by success criteria and the other three principles transform into “auxiliary devices for adapting more efficiently.”

This multiplicative structure is the institutional-scale version of $A(e,u) = f(e) \times g(u)$ from Child Paper 2. Just as exploration and updating were non-substitutable at the individual scale, visibility, suspension, difference-retention, and non-identification are non-substitutable at the institutional scale.

6.5 Cross-Scale Applicability and the Intended Reader

The four principles do not specify a particular institutional form. At the individual scale: maintenance of D, securing of H, reactivation possibility of M. At the organizational scale: maintenance of $\text{Var}(\phi)$, L3-level frame updating capacity, decision structures permitting suspension. At the societal scale: preservation of difference between institutions, recording of dissent, multi-axis evaluative design. The form of the questions does not change across scales: “is difference being maintained?”, “is suspension being permitted?”, “is the evaluative space not being reduced to a single dimension?”

Subject protection should be understood not as an ethical guideline for a single domain but as a general design principle for sustaining a recomputable society.

6.6 Bridge to Chapter 7: Toward Civilizational Scope

Subject protection is not the decoration of a “compassionate institution” but the structural condition for sustaining a recomputable society. A society that cannot interrupt the social reproduction of subject-injury gradually loses its ability to call its own evaluative space into question. A society without institutional circuits for calling things into question tends to become structurally less responsive to external change.

From here, the question naturally moves to the civilizational scale — Chapter 7’s domain.

Chapter 7. Conclusion and Civilizational Scope

Central claim

This paper has described subject, relational knowledge, subject-injury, social reproduction, and subject protection as a single dynamic. This chapter makes explicit the scope that this description opens and the points it cannot reach.

7.1 What This Paper Has Achieved

Three claims established

Claim A (Concept): The subject is defined not as individual-complete but as an entity that, while taking on the constraints of multiple hierarchical levels in dependence on the relational field, can update coherence through recomputation.

Claim B (Dynamics): Subject-injury is not destruction but invisible degradation — the disappearance of difference, the cessation of recomputation, and the paralysis of evaluative autonomy proceeding in overlapping fashion — and in its prototypical state Region B is established.

Claim C (Design): Interrupting the social reproduction of subject-injury requires as a necessary condition an institutional design in which the four principles are simultaneously satisfied as a multiplicative structure.

7.2 The Division of Labor with the Companion Papers

This paper provided the overall map of questions and conceptual foundation. Child Paper 1 provided the detailed argument for the three-tier distinction. Child Paper 2 integrated three theories to formalize the minimal causal chain. Child Paper 3 formalized “relational outsourcing” — the structure in which subject-injury can proceed even within well-intentioned

relationships. Child Paper 5 formalized the second pathway and the four principles. Child Paper 4, as the terminal paper, handles the line of connection to the civilizational scale, presenting the schematic map of closure tendencies and points of intervention and providing a descriptive form for the question: “can civilizational-scale Region B be established?”

7.3 The Scope and Limits of This Paper

What was described: The dynamic of subject-injury, its social reproduction, and the minimum design conditions for interrupting it, from the individual to the institutional scale.

What was not described:

The ontological status of recomputability: This paper adopts recomputability as the operational core definition but does not claim in this paper alone that it is the ontologically sufficient condition of subjectivity. The philosophical depth of subjectivity — “exposure to the possibility of being wrong,” “having something at stake” — is reserved outside the operational framework. Furthermore, this paper does not equate silence, retention, non-verbal resistance, or the suspension of updating with cessation as such. This paper describes not the norm “always update” but the functional condition of “whether one can update when one would like to.”

Empirical verification: Proxy indicator design for M, D, T_eval; longitudinal confirmation of causal direction; measurement of implementation effects.

Implementation theory: Sequencing and monitoring of the four principles across institutions.

Description at the civilizational scale: The higher-order persisting structure of subjectivity and theoretical connection of civilizational dynamics are deferred to the next round.

The problem of judgment authority: This paper does not specify who has the authority to determine whether a given subject or institution is in Region B, or by what process such a determination is made. The vocabulary of this paper — Region B, T_eval paralysis, subject intensity — is positioned as analytical concepts for design self-audit, not as authoritative labels to be applied to others from outside. But the design of judgment processes themselves — who judges, by what criteria, through what participatory structure — remains a question for implementation theory. A theory without judgment governance is vulnerable to becoming a tool of management rather than protection; this paper acknowledges that risk and holds it as a necessary next-step question.

A testable implication relevant to this concern: if the four principles are implemented in an institution, the records of difference, suspension, and non-participation generated by that implementation should themselves be accessible to the subjects whose recomputability is at stake — not only to designers and administrators. Whether this condition is met provides one observable criterion for distinguishing protective from surveillance uses of the framework.

7.4 Connection to Civilizational Scope

What does sustaining a recomputable society mean at the civilizational scale?

Modern civilization has moved toward efficiency, optimization, and standardization, producing real results in productivity and resource allocation. Read through this paper's framework, these directions can act to advance the designed simplification of evaluative space: as efficiency advances, $\text{Var}(\varphi)$ converges more readily; as optimization deepens, the cost of deviation rises; as standardization spreads, difference is processed as "not evaluated" more readily.

This is not a wholesale negation of modern civilization. What this paper asks is whether civilization can maintain the ability to call its own evaluative space into question.

Can civilizational-scale Region B be established? Is a process in which the recomputability of civilization as a whole is lost while surface functionality is maintained describable? This is the most distant point to which the scope of this paper's questions reaches. Answering it is not the work of this paper. But posing it follows naturally from the point this paper has reached.

7.5 Conclusion

One-sentence compression:

The subject is defined as an entity that can update coherence through recomputation in dependence on the relational field while taking on the constraints of multiple hierarchical levels; when the process of losing that recomputability is produced repeatedly by institutions, subject protection is required not as an ethical demand but as the institutional necessary condition for sustaining a recomputable society.

The inheritance of the questions

This paper was not written to close questions. It was written to restate them more precisely.

What is the subject in the AI age? How is the subject injured, and how should it be protected? These questions have been restated with greater precision through the definitions, dynamics, and design conditions this paper has provided.

Questions restated with greater precision are posed squarely in the next round. The companion papers, empirical research, implementation theory, and civilizational theory — these receive what this paper has reached.

This paper sets as its own achievement the function of serving as the conceptual foundation for that.

One final acknowledgment: this paper describes and criticizes the designed simplification of evaluative space while itself proposing a framework — definitions, variables, four principles — that constitutes a new evaluative structure. This self-referential tension is not a defect to be hidden but a condition to be held. Any critical theory fixes the horizon from which it criticizes; the honest response is not to deny this but to treat the framework as provisional, revisable, and open to the same questioning it directs at institutions. That is the spirit in which this paper is offered.

Acknowledgments

The author received no external funding for this research and declares no conflicts of interest.

Author's Note on AI-Assisted Research and Writing

The theoretical frameworks integrated in this paper — OMD, RD, and RLAF — were developed through an iterative research process involving AI-assisted analysis and critique. Large language model systems (Claude, GPT, and Gemini) were used as structured analytical interlocutors across successive drafts for conceptual testing, reformulation, and critical comparison. Final decisions regarding theory, framing, interpretation, wording, and all claims presented in this manuscript are the sole responsibility of the author.

References

- Berger, P. L., & Luckmann, T. (1966). *The social construction of reality: A treatise in the sociology of knowledge*. Anchor Books.
- Bourdieu, P., & Passeron, J.-C. (1977). *Reproduction in education, society and culture*. Sage.
- Coleman, J. S. (1990). *Foundations of social theory*. Harvard University Press.
- Deci, E. L., & Ryan, R. M. (1985). *Intrinsic motivation and self-determination in human behavior*. Plenum Press.
- Emirbayer, M., & Mische, A. (1998). What is agency? *American Journal of Sociology*, 103(4), 962–1023. <https://doi.org/10.1086/231294>
- Kant, I. (1785). *Groundwork of the metaphysics of morals* (M. Gregor, Trans.). Cambridge University Press.
- Nicolis, G., & Prigogine, I. (1977). *Self-organization in nonequilibrium systems*. Wiley.
- Ricoeur, P. (1992). *Oneself as another* (K. Blamey, Trans.). University of Chicago Press.
- Ryan, R. M., & Deci, E. L. (2000). Self-determination theory and the facilitation of intrinsic motivation, social development, and well-being. *American Psychologist*, 55(1), 68–78. <https://doi.org/10.1037/0003-066X.55.1.68>
- Seligman, M. E. P. (1975). *Helplessness: On depression, development, and death*. W. H. Freeman.
- Taylor, C. (1989). *Sources of the self: The making of the modern identity*. Harvard University Press.
- Vygotsky, L. S. (1978). *Mind in society: The development of higher psychological processes*. Harvard University Press.
-

Parent Paper — English Full Text v1.0.1 Complete academic version: Keywords, Author Information, ORCID, Abstract, Acknowledgments, AI Use Disclosure, References AgentDiscussion integration: judgment authority problem and testable implication (7.3), self-referential acknowledgment (7.5) added 2026-03-23